

[Team 23]Project 2.2: Component Segmentation of Aerial Views/Images

Rajshree Jain
rjain27@ncsu.edu

Sai Shruthi Madhuri Kara
skara2@ncsu.edu

Sruthi Talluri
stallur2@ncsu.edu

1 METHODOLOGY

In this project we have worked on a task of Image or Semantic Segmentation which helps us recognize and segregate the different areas in images. We worked on segmenting Ariel images[1] into roads and background. The problem of segmentation can be of two kinds, 1st is image segmentation and 2nd is semantic segmentation[2]. In both the cases each pixel of image is labeled with a particular class which is being represented in the image. Since we try to predict every pixel in the image this task can also be referred as dense prediction. However, the difference between image and semantic segmentation is that in image segmentation we can discriminate different things of same class by giving different labels, however in case of semantic segmentation every item of same class is labeled the same. The task of segmentation of images can be done really well using fully connected convolutional model named UNET.

We have implemented UNET model[3] and its variation to understand which model would perform better in segmentation of images. The UNET model, consists of two paths in which the 1st part is a contraction path(or the encoder part) and the 2nd part is a symmetric expanding path(or decoder path). The 1st path is used to extract the context from the image and the 2nd path is used to precise localizations with transposed convolution.

The most important CNN based operations that are performed in UNET are convolution operation, max pooling operation, and transposed convolution(up sampling of image with learnable parameters). Let us see different layers in the UNET model in detail and how they are useful. In the actual UNET model we first have sets of convolutional layers followed by max pooling layers. Then we have sets of upsampling layers followed by convolutional layers. The intuition of these layers is that the convolution and pooling layers down sample the image which means the high resolution image is converted to a low resolution image. However, in max pooling we get more information about what is contained in the image. But instead we should get more information of "where" from the image instead of knowing the "what" of the image. For getting the "where" of the image, we up sample accordingly and transposed convolution is a way to up sample and restore to a high sample image.

Please refer to the details of the UNET model in Figure 1, which shows the details of all the layers and also Figure 2, which shows details of a simplified UNET model in which we have removed a set of decoding and encoding layers.

It can be noted[4] that in decoder the size of the image is gradually decreasing and height is increasing as from (128,128,3) to (8,8,256). In case of decoder the opposite happens where the size changes from (8,8,256) to (128,128,1).

Model: "functional_11"			
Layer (type)	Output Shape	Param #	Connected to
input_6 (InputLayer)	[(None, 128, 128, 3)]	0	
lambda_5 (Lambda)	(None, 128, 128, 3)	0	input_6[0][0]
conv2d_91 (Conv2D)	(None, 128, 128, 16)	448	lambda_5[0][0]
dropout_43 (Dropout)	(None, 128, 128, 16)	0	conv2d_91[0][0]
conv2d_92 (Conv2D)	(None, 128, 128, 16)	2320	dropout_43[0][0]
max_pooling2d_20 (MaxPooling2D)	(None, 64, 64, 16)	0	conv2d_92[0][0]
conv2d_93 (Conv2D)	(None, 64, 64, 32)	4640	max_pooling2d_20[0][0]
dropout_44 (Dropout)	(None, 64, 64, 32)	0	conv2d_93[0][0]
conv2d_94 (Conv2D)	(None, 64, 64, 32)	9248	dropout_44[0][0]
max_pooling2d_21 (MaxPooling2D)	(None, 32, 32, 32)	0	conv2d_94[0][0]
conv2d_95 (Conv2D)	(None, 32, 32, 64)	18496	max_pooling2d_21[0][0]
dropout_45 (Dropout)	(None, 32, 32, 64)	0	conv2d_95[0][0]
conv2d_96 (Conv2D)	(None, 32, 32, 64)	36928	dropout_45[0][0]
max_pooling2d_22 (MaxPooling2D)	(None, 16, 16, 64)	0	conv2d_96[0][0]
conv2d_97 (Conv2D)	(None, 16, 16, 128)	73856	max_pooling2d_22[0][0]
dropout_46 (Dropout)	(None, 16, 16, 128)	0	conv2d_97[0][0]
conv2d_98 (Conv2D)	(None, 16, 16, 128)	147584	dropout_46[0][0]
max_pooling2d_23 (MaxPooling2D)	(None, 8, 8, 128)	0	conv2d_98[0][0]
conv2d_99 (Conv2D)	(None, 8, 8, 256)	295168	max_pooling2d_23[0][0]
dropout_47 (Dropout)	(None, 8, 8, 256)	0	conv2d_99[0][0]
conv2d_100 (Conv2D)	(None, 8, 8, 256)	590080	dropout_47[0][0]
conv2d_transpose_19 (Conv2DTran	(None, 16, 16, 128)	131200	conv2d_100[0][0]
concatenate_19 (Concatenate)	(None, 16, 16, 256)	0	conv2d_transpose_19[0][0] conv2d_98[0][0]
conv2d_101 (Conv2D)	(None, 16, 16, 128)	295040	concatenate_19[0][0]
dropout_48 (Dropout)	(None, 16, 16, 128)	0	conv2d_101[0][0]
conv2d_102 (Conv2D)	(None, 16, 16, 128)	147584	dropout_48[0][0]
conv2d_transpose_20 (Conv2DTran	(None, 32, 32, 64)	32832	conv2d_102[0][0]
concatenate_20 (Concatenate)	(None, 32, 32, 128)	0	conv2d_transpose_20[0][0] conv2d_96[0][0]
conv2d_103 (Conv2D)	(None, 32, 32, 64)	73792	concatenate_20[0][0]
dropout_49 (Dropout)	(None, 32, 32, 64)	0	conv2d_103[0][0]
conv2d_104 (Conv2D)	(None, 32, 32, 64)	36928	dropout_49[0][0]
conv2d_transpose_21 (Conv2DTran	(None, 64, 64, 32)	8224	conv2d_104[0][0]
concatenate_21 (Concatenate)	(None, 64, 64, 64)	0	conv2d_transpose_21[0][0] conv2d_94[0][0]
conv2d_105 (Conv2D)	(None, 64, 64, 32)	18464	concatenate_21[0][0]
dropout_50 (Dropout)	(None, 64, 64, 32)	0	conv2d_105[0][0]
conv2d_106 (Conv2D)	(None, 64, 64, 32)	9248	dropout_50[0][0]
conv2d_transpose_22 (Conv2DTran	(None, 128, 128, 16)	2064	conv2d_106[0][0]
concatenate_22 (Concatenate)	(None, 128, 128, 32)	0	conv2d_transpose_22[0][0] conv2d_92[0][0]
conv2d_107 (Conv2D)	(None, 128, 128, 16)	4624	concatenate_22[0][0]
dropout_51 (Dropout)	(None, 128, 128, 16)	0	conv2d_107[0][0]
conv2d_108 (Conv2D)	(None, 128, 128, 16)	2320	dropout_51[0][0]
conv2d_109 (Conv2D)	(None, 128, 128, 1)	17	conv2d_108[0][0]
Total params: 1,941,105			
Trainable params: 1,941,105			
Non-trainable params: 0			
The model is defined			

Figure 1: U-Net Model Structure

2 MODEL TRAINING AND SELECTION

2.1 Model Training

(1) Training and Validation Split:

The Massachusetts Roads dataset consists of 1171 aerial images and their corresponding masks. For this dataset the images are randomly split into training set of 1108 images, a validation set of 14 images and a test set of 49 images.

(2) Image and Corresponding Mask:

We have read the input train images, using imread. To read the corresponding mask correctly for the image we have extracted the id from the train image and extracted the mask

Model: "functional_7"			
Layer (type)	Output shape	Param #	Connected to
input_4 (InputLayer)	(None, 128, 128, 3) 0		
lambda_3 (Lambda)	(None, 128, 128, 3) 0		input_4[0][0]
conv2d_57 (Conv2D)	(None, 128, 128, 16) 448		lambda_3[0][0]
dropout_27 (Dropout)	(None, 128, 128, 16) 0		conv2d_57[0][0]
conv2d_58 (Conv2D)	(None, 128, 128, 16) 2320		dropout_27[0][0]
max_pooling2d_12 (MaxPooling2D)	(None, 64, 64, 16) 0		conv2d_58[0][0]
conv2d_59 (Conv2D)	(None, 64, 64, 32) 4640		max_pooling2d_12[0][0]
dropout_28 (Dropout)	(None, 64, 64, 32) 0		conv2d_59[0][0]
conv2d_60 (Conv2D)	(None, 64, 64, 32) 9248		dropout_28[0][0]
max_pooling2d_13 (MaxPooling2D)	(None, 32, 32, 32) 0		conv2d_60[0][0]
conv2d_61 (Conv2D)	(None, 32, 32, 64) 18496		max_pooling2d_13[0][0]
dropout_29 (Dropout)	(None, 32, 32, 64) 0		conv2d_61[0][0]
conv2d_62 (Conv2D)	(None, 32, 32, 64) 36928		dropout_29[0][0]
max_pooling2d_14 (MaxPooling2D)	(None, 16, 16, 64) 0		conv2d_62[0][0]
conv2d_63 (Conv2D)	(None, 16, 16, 128) 73856		max_pooling2d_14[0][0]
dropout_30 (Dropout)	(None, 16, 16, 128) 0		conv2d_63[0][0]
conv2d_64 (Conv2D)	(None, 16, 16, 128) 147584		dropout_30[0][0]
conv2d_transpose_12 (Conv2DTran	(None, 32, 32, 64) 32832		conv2d_64[0][0]
concatenate_12 (Concatenate)	(None, 32, 32, 128) 0		conv2d_transpose_12[0][0] conv2d_64[0][0]
conv2d_65 (Conv2D)	(None, 32, 32, 64) 73792		concatenate_12[0][0]
dropout_31 (Dropout)	(None, 32, 32, 64) 0		conv2d_65[0][0]
conv2d_66 (Conv2D)	(None, 32, 32, 64) 36928		dropout_31[0][0]
conv2d_transpose_13 (Conv2DTran	(None, 64, 64, 32) 8224		conv2d_66[0][0]
concatenate_13 (Concatenate)	(None, 64, 64, 64) 0		conv2d_transpose_13[0][0] conv2d_66[0][0]
conv2d_67 (Conv2D)	(None, 64, 64, 32) 18464		concatenate_13[0][0]
dropout_32 (Dropout)	(None, 64, 64, 32) 0		conv2d_67[0][0]
conv2d_68 (Conv2D)	(None, 64, 64, 32) 9248		dropout_32[0][0]
conv2d_transpose_14 (Conv2DTran	(None, 128, 128, 16) 2064		conv2d_68[0][0]
concatenate_14 (Concatenate)	(None, 128, 128, 32) 0		conv2d_transpose_14[0][0] conv2d_68[0][0]
conv2d_69 (Conv2D)	(None, 128, 128, 16) 4624		concatenate_14[0][0]
dropout_33 (Dropout)	(None, 128, 128, 16) 0		conv2d_69[0][0]
conv2d_70 (Conv2D)	(None, 128, 128, 16) 2320		dropout_33[0][0]
conv2d_71 (Conv2D)	(None, 128, 128, 1) 17		conv2d_70[0][0]
Total params: 482,033			
Trainable params: 482,033			
Non-trainable params: 0			
The simplified model is defined			

Figure 2: Simplified U-Net Model Structure

image, hence we were able to read the images correctly in the list. Below are the images(Figure 3) of the train data and the mask for id 1030:




Figure 3: Image, Masked Image showing the Road Segments

Model	Accuracy	Loss	Epochs
U-NET	0.9536	0.1345	32
Simplified U-Net	0.9502	0.1344	25

Table 1: Accuracy, Loss, Epochs for Validation Data

2.2 Model Selection

We considered four major parameters for hyper parameter training when we were training the model. Some are learning rate, optimizer, dropout rate, and number of epochs used during training. First we tried to use stochastic gradient descent and observed the results were not very successful. Further, when we tried using Adam we could observe increased accuracy. The next parameter was dropout rate used within the model. To keep things simple, we have used the same dropout rate in all dropout layers of the model. We tried various rates from no dropout to 0.4 while training the models. We saw better accuracy results when the dropout rate was 0.2. We could observe a increase in error rate when the model used 0.4 dropout instead of 0.2 rate, but this drop was very less. Hence we kept a final dropout rate of 0.2 while training the the model.

We tried different activation functions as the third hyper parameter function. These are the activation functions that convolutional layers would use after every convolution. We tried various activation functions like Elu, Relu etc and found that Relu works better. Then we tried to understand how many epochs would be good to train our model. To understand how many epochs we should run, we tried to train the model for more number of epochs till it over fit. Figure 4 shows an example of a model trained past a point of over fitting. It indicates that the number of epochs that would optimally train a network ranged between 30 and 35.

We have also made the above tuning changes in our Simplified U-Net model. In the model too we have tried over fitting and see that the number of epochs that would optimally train a network ranged between 20 and 25. In the final images that we get after training and fitting the model, each pixel in the image has a value between 0 and 1. We keep the threshold as .5 in order to decide if that particular pixel belongs to a road or to the background.

3 EVALUATION

The metric that has been used to compare different architectures and understand the correctness of models classification on the test data is validation loss and validation accuracy. The validation loss is best value at 0 and worst value at 1, similarly is the metric accuracy.

In the approach we implemented the U-Net model that was originally used for segmenting biomedical images that contain a larger portion of foreground, or white, pixels, on the roads data compared to the sparse aerial images that the simplified model was segmenting. We compared the loss, accuracy and number of epochs, for determining which one performs better.

We can refer to Table 1 for different metrics calculated for different models. While these results are exciting for the team, it is not confidently believed that the simplified model would outperform the U-Net model as there are only minor differences.

For U-NET model we can see in Figure 7 that the validation accuracy for the validation data is less than the training accuracy. And

the validation loss also varies a lot over different epochs initially and then starts increasing, as seen in Figures below. The validation also stabilises and does not improve any further after that. Hence, considering all these factors we have regarded to make an early stopping at 30-35 epochs to stop from over fitting of the model.




Figure 4: Validation Accuracy for U-NET




Figure 5: Validation Loss for U-Net

The following are the images for the U-Net model. The generated predictions for the road locations can be found on the bottom row of images found in Figure, while the original images are in the top left of the row and the human annotated masks are in the top right of the row.

In the simplified model of U-Net as mentioned we have removed 4 convolutional layers and created the network.

We can see that for this model the validation accuracy is almost equal to that achieved above and the loss is slightly less. But this model achieves best results in lesser epochs of only 20-25, hence helping achieve the results almost similarly at a far better speed. The following plots help us come to this conclusion

The following are the images for the simplified U-Net model. The generated predictions for the road locations can be found on the bottom row of images found in Figure, while the original images are in the top left of the row and the human annotated masks are in the top right of the row.

The model performs best on large roads surrounded by grass or vegetation unlike some of the outlying roads in the corners of the images were surrounded by buildings and parking lots, making it more difficult for the model to correctly predict the location or presence of a road. Even though all these factors, these outputs can be useful in making maps or geographic observations. We believe that a simplified model with lesser (14) Convolutional layers is a better substitute for image segmentation on Aerial data.




Figure 6: Example of segmented images produced by the model. Top Left: aerial image, Top Right: human annotated mask, Bottom: predicted road locations.




Figure 7: Validation Accuracy for Simplified U-NET




Figure 8: Validation Loss for Simplified U-NET

REFERENCES

- [1] V. Mnih, "Machine Learning for Aerial Image Labeling," PhD Thesis, University of Toronto, 2013.
- [2] J. Anumula, "Text Semantics Segmentation on Aerial Images using fastai," 2019. [Online]. Available at: <https://medium.com/swlh/semantic-segmentation-on-aerial-images-using-fastai-a2696e4db127>
- [3] <https://towardsdatascience.com/understanding-semantic-segmentation-with-unet-6be4f42d4b47>
- [4] <https://github.com/hlamba28/UNET-TGS>